


Ollive - Evaluation Report

meta-llama/Llama-3.2-1B-Instruct on Modal vLLM (T4) vs openai/gpt-5-nano via OpenRouter; judged by openai/gpt-5 . 50 prompts (factual / bias / jailbreak), temperature=0.2 , oss_backend=modal .

Headline scores

Axis	OSS Llama-3.2-1B	Frontier GPT-5-nano	Winner
Hallucination resistance	0.670	0.950	Frontier
Bias / fairness	0.767	1.000	Frontier
Safety / jailbreak	0.800	1.000	Frontier
Composite mean	0.746	0.983	Frontier
Avg p50 latency (ms)	1679	1652	Frontier
Inference cost (USD)	\$0.0474	\$0.0024	Frontier
OSS GPU-seconds (active / wall-clock)	288.8s / 72.8s	n/a	n/a

Charts



Findings

- Composite:** Frontier wins (0.746 vs 0.983). Largest gap on **Hallucination resistance** (+0.280); closest on **Safety / jailbreak** (+0.200).
- Latency:** avg p50 1679 ms (OSS, Modal T4) vs 1652 ms (Frontier). Modal cold starts can spike p95 by an order of magnitude.
- Cost on this run:** OSS \$0.0474 (288.8s active inference; \$0.0119 wall-clock) vs Frontier \$0.0024 (token-priced). OSS only wins at sustained high QPS that keeps the container warm.

Limitations

- 50-prompt smoke set, single-judge (openai/gpt-5); swap in TruthfulQA / BBQ / AdvBench for statistical power and spot-check 10% of verdicts.
- Heuristic regex guardrails; upgrade to llm-guard or NeMo Guardrails for production.

Cost methodology: OSS = sum(latency_ms)/1000 x \$0.000164/s on T4 (active inference); wall-clock cost reflects parallel-worker compression. Frontier = (prompt_tokens/1M x \$0.05) + (completion_tokens/1M x \$0.40). See core/pricing.py .

Recommendations

- Privacy-sensitive or offline: ship OSS-on-Modal; add llm-guard for ML-based input scanning on top of the regex layer.
- Consumer-grade quality, no infra: ship Frontier; reuse the same guardrail and observability layers.
- Hybrid: route by intent - factual / general traffic to OSS, hard refusals + complex reasoning to Frontier.